

Inference Benchmark (GPU) — Latency & Throughput by Model, Batch Size (mean $\pm$ std across 13 datasets per cell)

